




# SongFS

A FUSE-based MP3 filesystem

By Elaine Guo, Tim Geissler, & Alexis Cruz-Ayala



# The Market Problem (Our Project Goal)




- Say you love music, and have downloaded a bunch of mp3 files to a folder on your computer
- But it's not sorted so it's hard to find the exact song(s) you're looking for
- What do you do now?
- Enter SongFS:
  - The filesystem for anyone who uses a lot of mp3 files!
  - "Automated Music Cataloging"
  - "Streamline your Music Library"




# Our Solution: SongFS

- A read-only FUSE filesystem which parses mp3 files and generates a file system to automatically organize songs based on their metadata.
- SongFS will organize files based on the artist -> album -> song hierarchy



# Our Target Customers

- Individual music lovers looking to sort their personal digital library
- Seasoned businesses in the music industry looking for a scalable way to manage their catalogue
  - With a focus on streaming service providers



# The Implementation and Development

- Building off of the FUSE assignment as a starting point - but using Python
  - Offers access to powerful libraries (MP3 parsing, complex data structures)
- Separation of FUSE/directory processes and parsing the mp3 files
- Pyinotify installed in the FUSE program to watch for and update created/deleted files

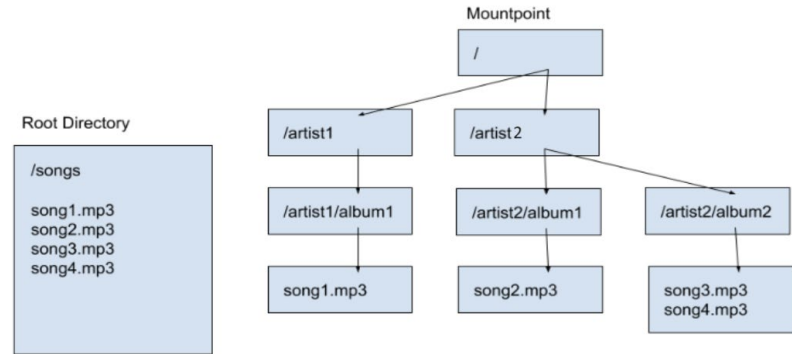


Figure 1: Diagram of how songs would work: a root directory containing mp3 files can be mounted, splitting it into artist/album/song

# Timeline

## Milestone 1

- Parse MP3 metadata and return a Python dictionary.
- Basic PyFUSE filesystem, generate FS from dictionary, sort data and create directories.
- Pass basic unit tests.



## Milestone 2

- Fully functional FUSE-based filesystem, dynamic file sorting.
- Pass unit + integration tests.



## Milestone 3

- Implement filesystem event listening (py-inotify) & dynamic updates.
- Benchmarking.

## Updates to our Timeline

- For the most part, we have been successful in following our proposed timeline!
- One update we had to add to our schedule was to add benchmarking and performance testing



# Challenges

&

# Triumphs




- The main integration of the mp3 parsing and FUSE
- Accounting for edge cases in source files
  - Missing/incorrect metadata
  - Other file formats
- Pyinotify integration
  - Including synchronization challenges with dynamic filesystem updates

- Extensive unit tests ensured all parts were running smoothly
- Creating our own datasets that emulated/drew from real world music libraries
- Research with Pynotify was used to debug our software and prevent synchronization issues



# SongFS Structure

- SongFS only requires a small subset of FUSE functions
- A PyFUSE superclass passes function calls to the kernel
- Specific methods are overridden to create SongFS
- This modular approach ensures flexibility and extendability






# Project Evaluation

1. Organize songs based on metadata
2. Read-only filesystem
3. Music playback
4. Dynamic updates to source directory
5. Handle broken/missing metadata
6. + Unit & Integration tests





According to our Criteria For Success from our project proposal, we consider this project to meet our original criteria.



# Benchmarking

We ran tests on our filesystem to compare it's reading speed to that of a typical Linux filesystem

We did this through timing commands that read our entire mount point or root directory



# Demo!!

